Introduction to Agents

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Agenda

1. Introduction to Agents
2. Core Components of Agent Architecture
3. Difference Between LLM Workflows and Agents
4. Single Agent Design Patterns
5. ReAct Framework for Planning Agents
6. Agent Memory, Self-Reflecting Agents, and MCP

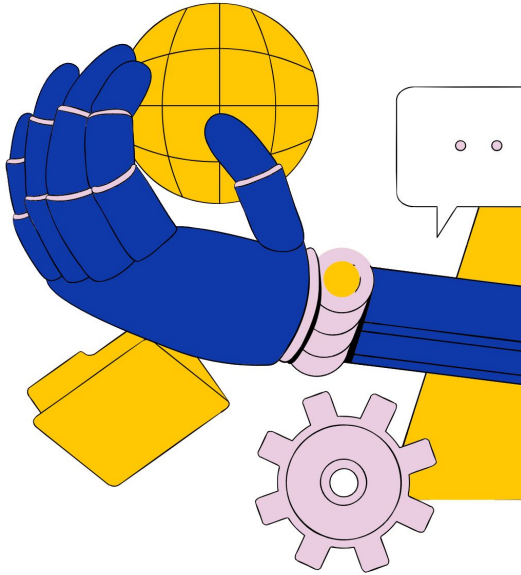

You Say It. An AI Agent Does It.

> "Book me a flight to Delhi tomorrow morning."

What happens next is remarkable. The system doesn't just reply with information — it **takes action**. It understands your request, searches available flights, compares options, and books the best one. This is the difference between a chatbot that answers and an AI agent that **gets things done**.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Made with GAMMA

What are Agents?

- AI agents are **systems that actively engage in tasks**, often involving planning, decision-making, and interaction with their environment to achieve defined objectives.

- AI Agents can perceive their environment and act upon that environment.

- In contrast, systems built around traditional AI models passively provide outputs in response to specific inputs.

- Example:
  Self-driving cars serve as a **prime example of AI agents**. They utilize **sensors**, **cameras**, and **radars** to perceive their surroundings. The system processes this environmental data and employs actuators to take action, making decisions based on its **cognitive modules**. Over time, all of these **modules learn to adapt to new patterns**, continuously refining their driving skills accordingly.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

What are Agents?

AI Agents: Core Components

Agents are goal-oriented, autonomous systems composed of distinct functional components.

Agents

Planning

Capacity to set goals, break down complex tasks into sub-tasks, and devise strategies or sequences of actions to achieve objectives.

Reasoning

The ability to process information, understand context, make inferences, and formulate decisions.

Execution

The ability to interact with the environment and carry out the planned actions through tool execution.

Memory

Mechanisms for storing and retrieving past experiences, knowledge, and context.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

What Exactly Is an AI Agent?

An **AI Agent** is a system that can understand, decide, and act to complete a task — like a smart assistant that doesn't just respond, it **works for you**.

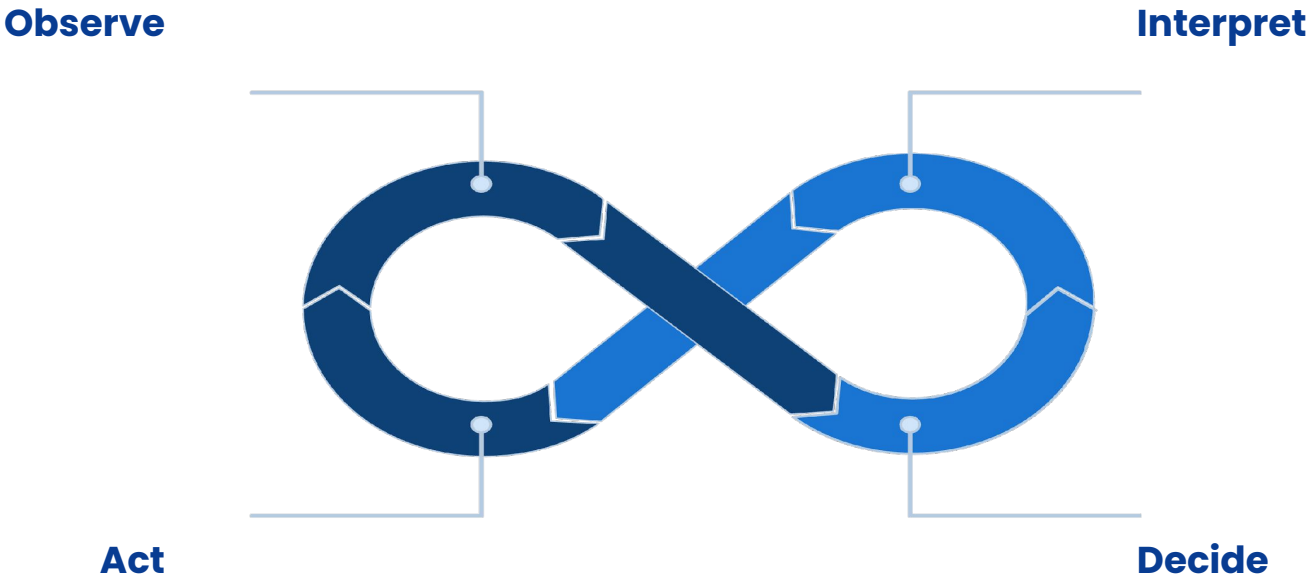

This observe–think–act loop is what separates AI agents from simple chatbots. They don't just generate text — they execute real-world tasks on your behalf.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Made with GAMMA

Where Do AI Agents Show Up?

AI agents are already working behind the scenes in everyday tools and services.

Customer Support

Resolve queries, process refunds, and escalate issues — instantly, without human wait times.

Smart Assistants

Set reminders, control devices, manage schedules, and handle daily tasks on command.

Shopping Assistants

Find products, compare prices, and complete purchases tailored to your preferences.

Banking Assistants

Track spending, flag fraud, and automate transfers — your personal financial co-pilot.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Made with GAMMA

Why AI Agents Are a Big Deal

⏱️ Save Time

Automate repetitive work so teams focus on high-value decisions.

⚡ Faster Decisions

Process data and act in seconds, not hours.

📈 Boost Productivity

Handle multiple tasks simultaneously, at scale.

🤝 Always Available

No breaks, no delays — support 24/7, 365 days a year.

The Future of Work

The most powerful organisations won't choose between humans and AI agents — they'll combine both.

Humans bring creativity and judgment. Agents bring speed and scale. Together, they unlock a new level of performance.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Made with GAMMA

Generative AI Agents

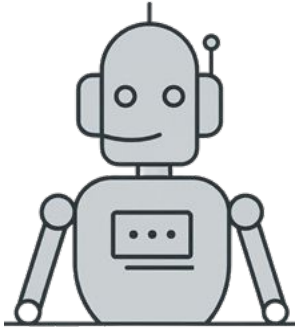

Generative AI agents are an evolving subset of AI Agents that use generative models to perceive their environment, produce novel and relevant outputs, make decisions, and interact with their surroundings to achieve goals with minimal human intervention.

Autonomy – Do not rely on pre-programmed rules or step-by-step instructions and can independently handle complex scenarios relying on reasoning afforded by LLMs.

Perception – Use text-based or multimodal inputs and interactions with other systems or users to perceive their environment.

Action – Determine the best course of action and execute tasks to fulfil their goals. Task execution is achieved by selecting, sequencing and calling a set of tools available to the agent.

Goal-Oriented Behavior – Designed to achieve specific, pre-defined goals by autonomously breaking down complex tasks into smaller steps.

**For the rest of this course, we use the terms agents and Generative AI agents interchangeably*

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Generative AI Agents

Agent = LLM (for planning/reasoning) + Tool Execution (for action)

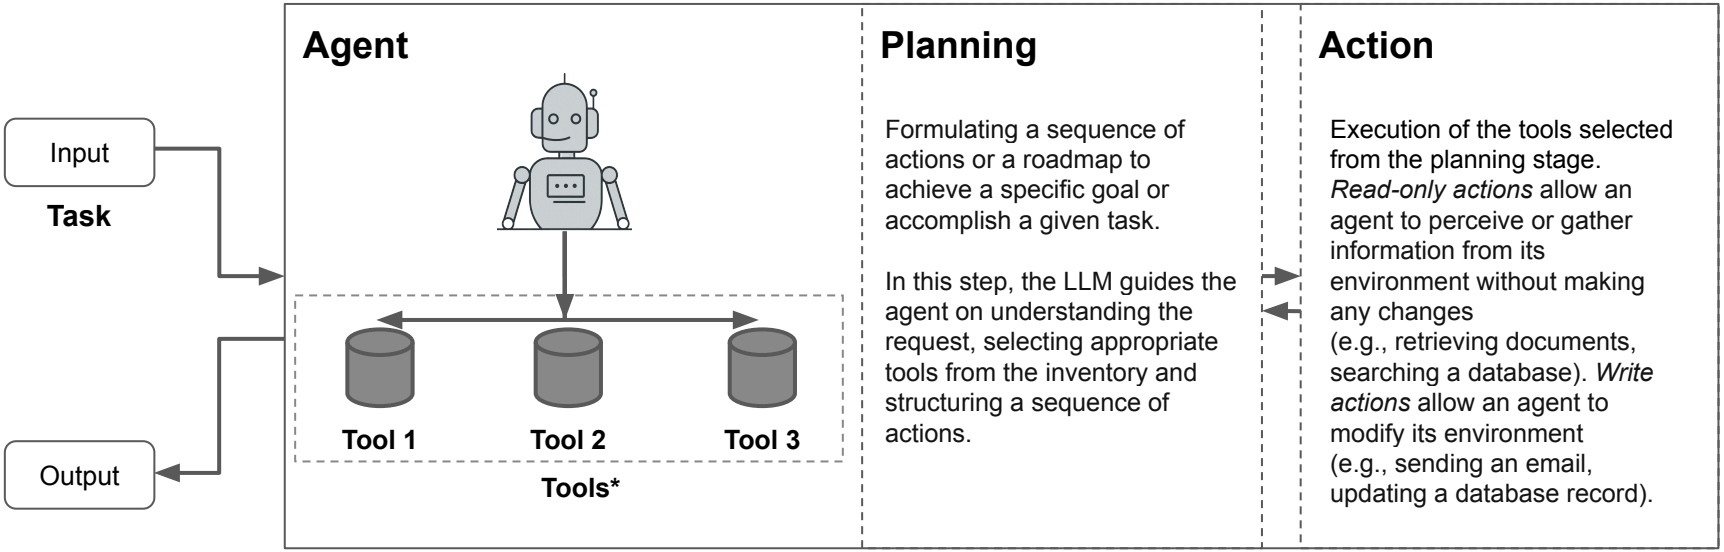

Tools are the specific resources (e.g., APIs, databases) the agent uses to find information or take action beyond its own knowledge. The agent evaluates the output from each tool step, replanning or escalating upon failure. Once successful, it uses the gathered information or completed actions to generate a natural language customer response.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Generative AI Agents: Use Case

Handling E-Commerce Customer Inquiries

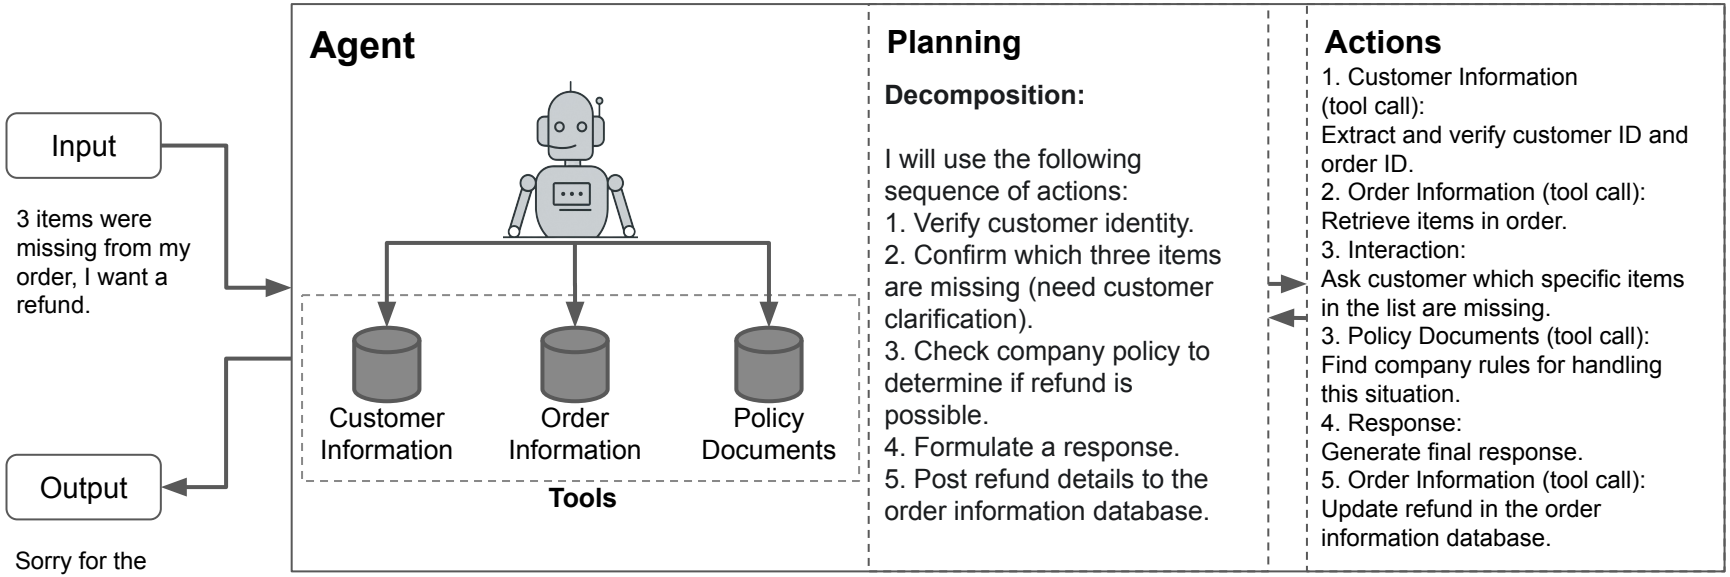

The agent looks at the entire trace of actions and their outputs and composes the final answer. If resolution is not possible within the action sequence, the query is escalated to a human.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Where Do AI Agents Show Up?

AI agents are already working behind the scenes in everyday tools and services.

Customer Support

Resolve queries, process refunds, and escalate issues — instantly, without human wait times.

Smart Assistants

Set reminders, control devices, manage schedules, and handle daily tasks on command.

Shopping Assistants

Find products, compare prices, and complete purchases tailored to your preferences.

Banking Assistants

Track spending, flag fraud, and automate transfers — your personal financial co-pilot.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Made with GAMMA

Why AI Agents Are a Big Deal

⏱️ Save Time

Automate repetitive work so teams focus on high-value decisions.

⚡ Faster Decisions

Process data and act in seconds, not hours.

📈 Boost Productivity

Handle multiple tasks simultaneously, at scale.

🤝 Always Available

No breaks, no delays — support 24/7, 365 days a year.

The Future of Work

The most powerful organisations won't choose between humans and AI agents — they'll combine both.

Humans bring creativity and judgment. Agents bring speed and scale. Together, they unlock a new level of performance.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Made with GAMMA

Generative AI Agents

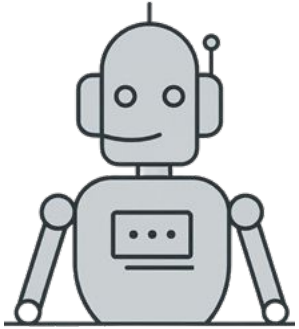

Generative AI agents are an evolving subset of AI Agents that use generative models to perceive their environment, produce novel and relevant outputs, make decisions, and interact with their surroundings to achieve goals with minimal human intervention.

Autonomy – Do not rely on pre-programmed rules or step-by-step instructions and can independently handle complex scenarios relying on reasoning afforded by LLMs.

Perception – Use text-based or multimodal inputs and interactions with other systems or users to perceive their environment.

Action – Determine the best course of action and execute tasks to fulfil their goals. Task execution is achieved by selecting, sequencing and calling a set of tools available to the agent.

Goal-Oriented Behavior – Designed to achieve specific, pre-defined goals by autonomously breaking down complex tasks into smaller steps.

**For the rest of this course, we use the terms agents and Generative AI agents interchangeably*

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Generative AI Agents

Agent = LLM (for planning/reasoning) + Tool Execution (for action)

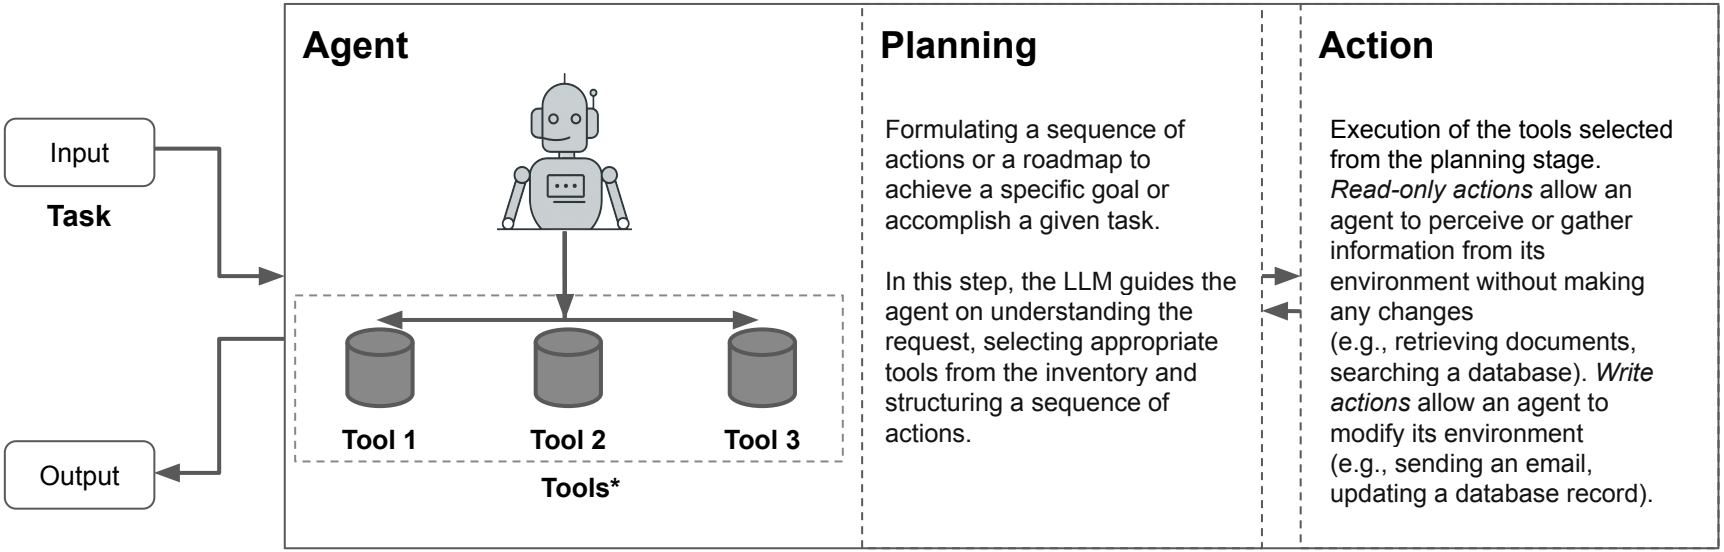

Tools are the specific resources (e.g., APIs, databases) the agent uses to find information or take action beyond its own knowledge. The agent evaluates the output from each tool step, replanning or escalating upon failure. Once successful, it uses the gathered information or completed actions to generate a natural language customer response.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Generative AI Agents: Use Case

Handling E-Commerce Customer Inquiries

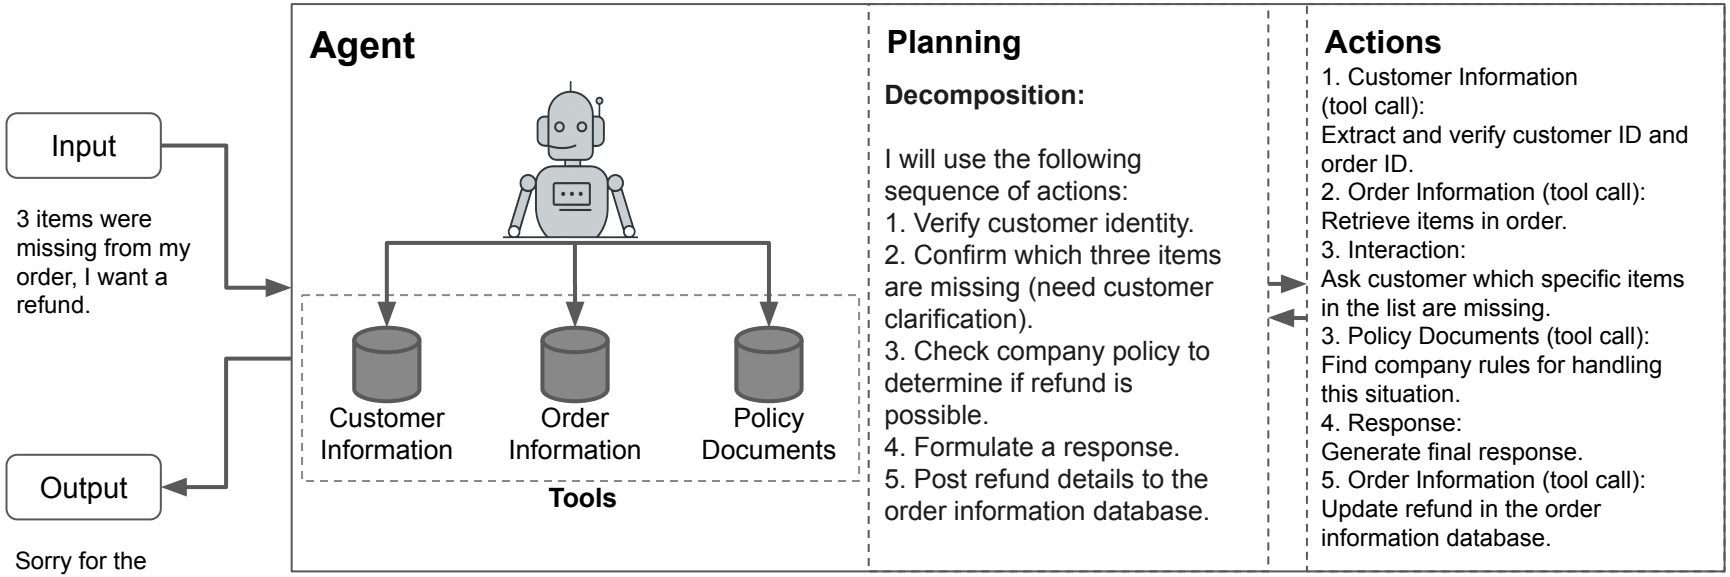

The agent looks at the entire trace of actions and their outputs and composes the final answer. If resolution is not possible within the action sequence, the query is escalated to a human.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Contextualizing GenAI Agents

	LLMs	GenAI Assistants	GenAI Workflows	GenAI Agents
What?	Advanced Machine Learning models that can perform singular natural language tasks.	LLMs that take on specific roles in an interactive manner. Maintain history of previous interactions.	Tasks are broken down into a fixed sequence of steps. Each step is executed by an LLM call (and possibly, tools).	Autonomously break down complex tasks into subtasks and carry them out using the tools at their disposal. Goal-oriented systems.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Contextualizing GenAI Agents

	LLMs	GenAI Assistants	GenAI Workflows	GenAI Agents
When to use?	Excel in tasks like summarization, text generation, translation, and creative writing. To provide coding assistance. Foundational block for applications involving reasoning.	Useful for customer support and service automation. Leverage external data sources for context using RAG. Foundational block for multi-step GenAI workflows.	Ideal for situations where the task can be easily and cleanly decomposed into fixed subtasks (e.g., generating marketing copy). Routing and parallelization can increase effectiveness of independent subtasks (e.g., specialized processes for customer queries).	Ideal for situations where it's difficult or impossible to predict the required number of steps and where a fixed path cannot be hardcoded. Goal-based agents excel in complex decision-making environments, evaluating potential actions, such as in financial trading.

For many organizations, starting with prescriptive GenAI workflows provides a more manageable and often more effective entry point than immediately deploying autonomous agents.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Internal Workings of A Single-agent System

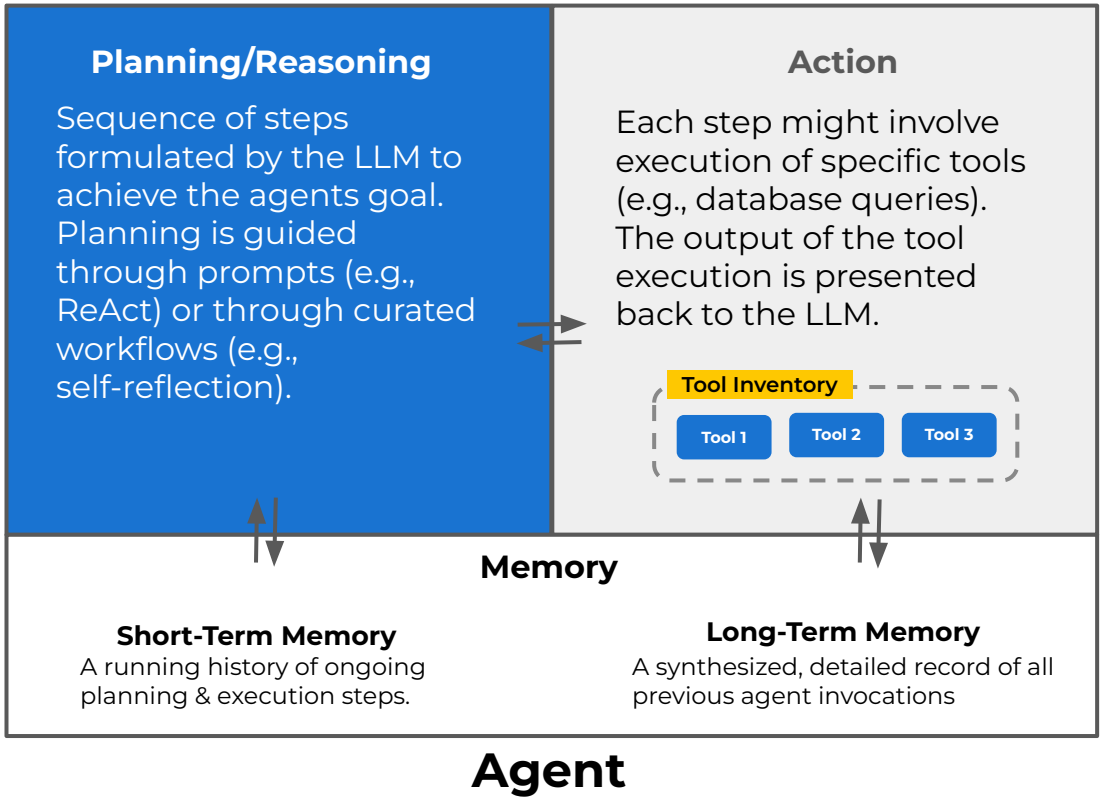

- A single-agent system operates through iterative cycles of planning, action, and memory.
- The LLM formulates a sequence of steps (Planning/Reasoning) to achieve a goal, guided by prompts or curated workflows. Each step may involve the LLM selecting and executing specific tools (Action).
- The results of these tool executions are then fed back into the LLM.
- The agent utilizes both short-term memory to track the current task and long-term memory to synthesize information from previous interactions.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Where Do AI Agents Show Up?

AI agents are already working behind the scenes in everyday tools and services.

Customer Support

Resolve queries, process refunds, and escalate issues — instantly, without human wait times.

Smart Assistants

Set reminders, control devices, manage schedules, and handle daily tasks on command.

Shopping Assistants

Find products, compare prices, and complete purchases tailored to your preferences.

Banking Assistants

Track spending, flag fraud, and automate transfers — your personal financial co-pilot.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Made with GAMMA

Why AI Agents Are a Big Deal

⏱️ Save Time

Automate repetitive work so teams focus on high-value decisions.

⚡ Faster Decisions

Process data and act in seconds, not hours.

📈 Boost Productivity

Handle multiple tasks simultaneously, at scale.

🤝 Always Available

No breaks, no delays — support 24/7, 365 days a year.

The Future of Work

The most powerful organisations won't choose between humans and AI agents — they'll combine both.

Humans bring creativity and judgment. Agents bring speed and scale. Together, they unlock a new level of performance.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Made with GAMMA

Generative AI Agents

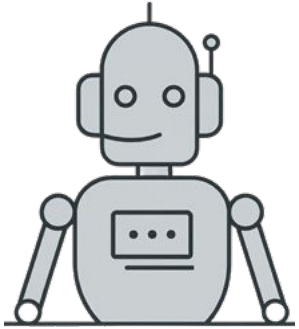

Generative AI agents are an evolving subset of AI Agents that use generative models to perceive their environment, produce novel and relevant outputs, make decisions, and interact with their surroundings to achieve goals with minimal human intervention.

Autonomy – Do not rely on pre-programmed rules or step-by-step instructions and can independently handle complex scenarios relying on reasoning afforded by LLMs.

Perception – Use text-based or multimodal inputs and interactions with other systems or users to perceive their environment.

Action – Determine the best course of action and execute tasks to fulfil their goals. Task execution is achieved by selecting, sequencing and calling a set of tools available to the agent.

Goal-Oriented Behavior – Designed to achieve specific, pre-defined goals by autonomously breaking down complex tasks into smaller steps.

**For the rest of this course, we use the terms agents and Generative AI agents interchangeably*

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Generative AI Agents

Agent = LLM (for planning/reasoning) + Tool Execution (for action)

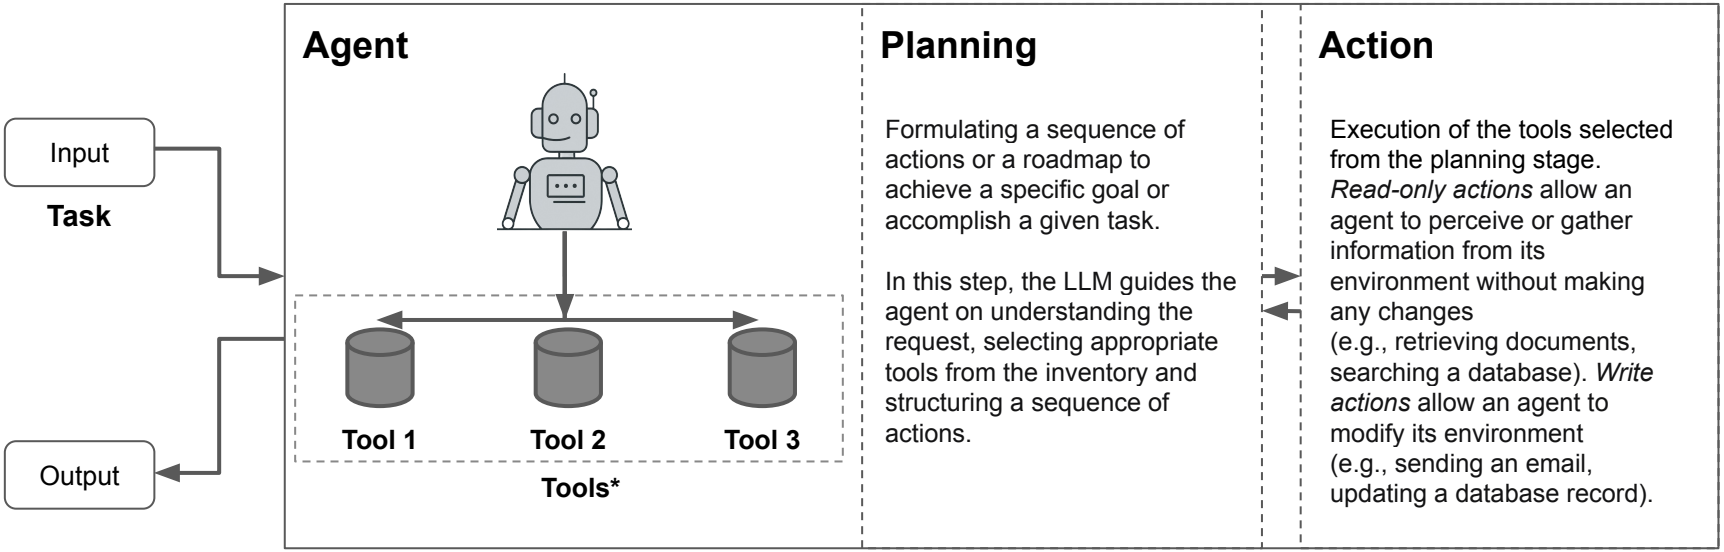

Tools are the specific resources (e.g., APIs, databases) the agent uses to find information or take action beyond its own knowledge. The agent evaluates the output from each tool step, replanning or escalating upon failure. Once successful, it uses the gathered information or completed actions to generate a natural language customer response.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Generative AI Agents: Use Case

Handling E-Commerce Customer Inquiries

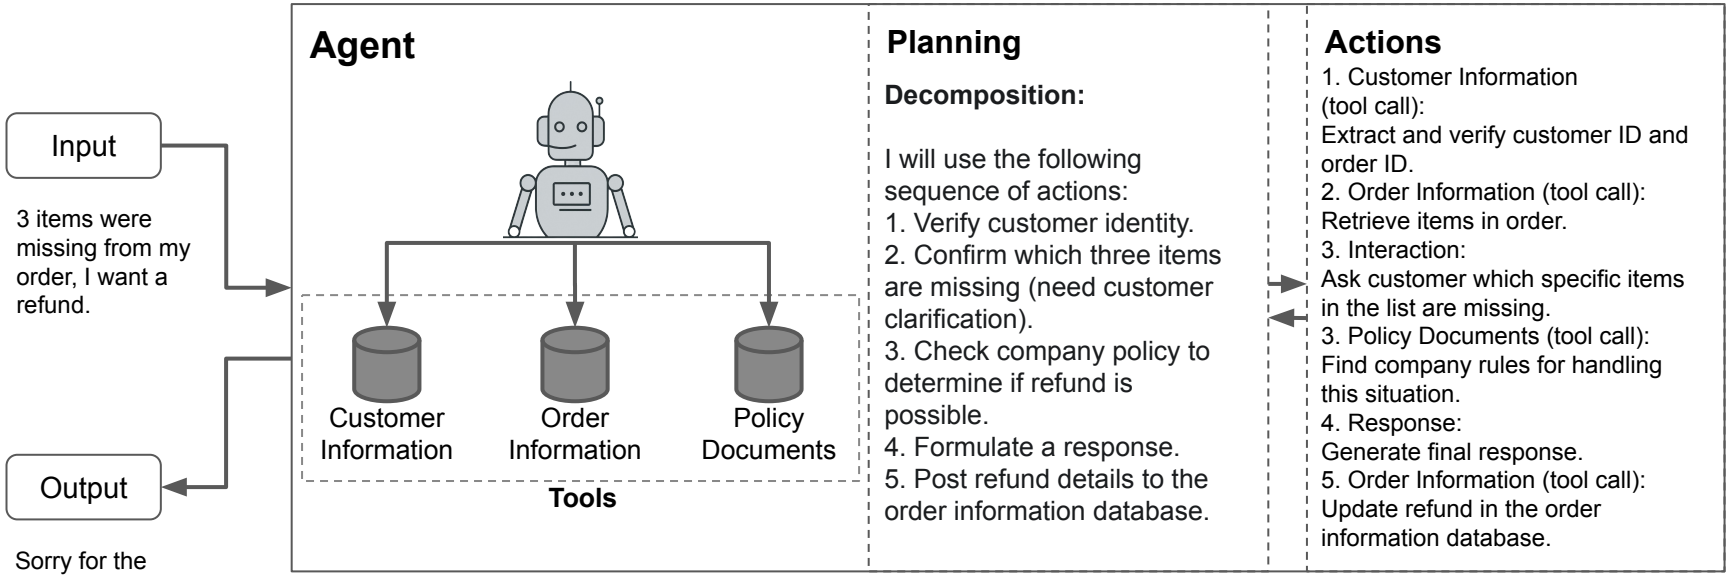

The agent looks at the entire trace of actions and their outputs and composes the final answer. If resolution is not possible within the action sequence, the query is escalated to a human.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Contextualizing GenAI Agents

	LLMs	GenAI Assistants	GenAI Workflows	GenAI Agents
What?	Advanced Machine Learning models that can perform singular natural language tasks.	LLMs that take on specific roles in an interactive manner. Maintain history of previous interactions.	Tasks are broken down into a fixed sequence of steps. Each step is executed by an LLM call (and possibly, tools).	Autonomously break down complex tasks into subtasks and carry them out using the tools at their disposal. Goal-oriented systems.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Contextualizing GenAI Agents

	LLMs	GenAI Assistants	GenAI Workflows	GenAI Agents
When to use?	Excel in tasks like summarization, text generation, translation, and creative writing. To provide coding assistance. Foundational block for applications involving reasoning.	Useful for customer support and service automation. Leverage external data sources for context using RAG. Foundational block for multi-step GenAI workflows.	Ideal for situations where the task can be easily and cleanly decomposed into fixed subtasks (e.g., generating marketing copy). Routing and parallelization can increase effectiveness of independent subtasks (e.g., specialized processes for customer queries).	Ideal for situations where it's difficult or impossible to predict the required number of steps and where a fixed path cannot be hardcoded. Goal-based agents excel in complex decision-making environments, evaluating potential actions, such as in financial trading.

For many organizations, starting with prescriptive GenAI workflows provides a more manageable and often more effective entry point than immediately deploying autonomous agents.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Internal Workings of A Single-agent System

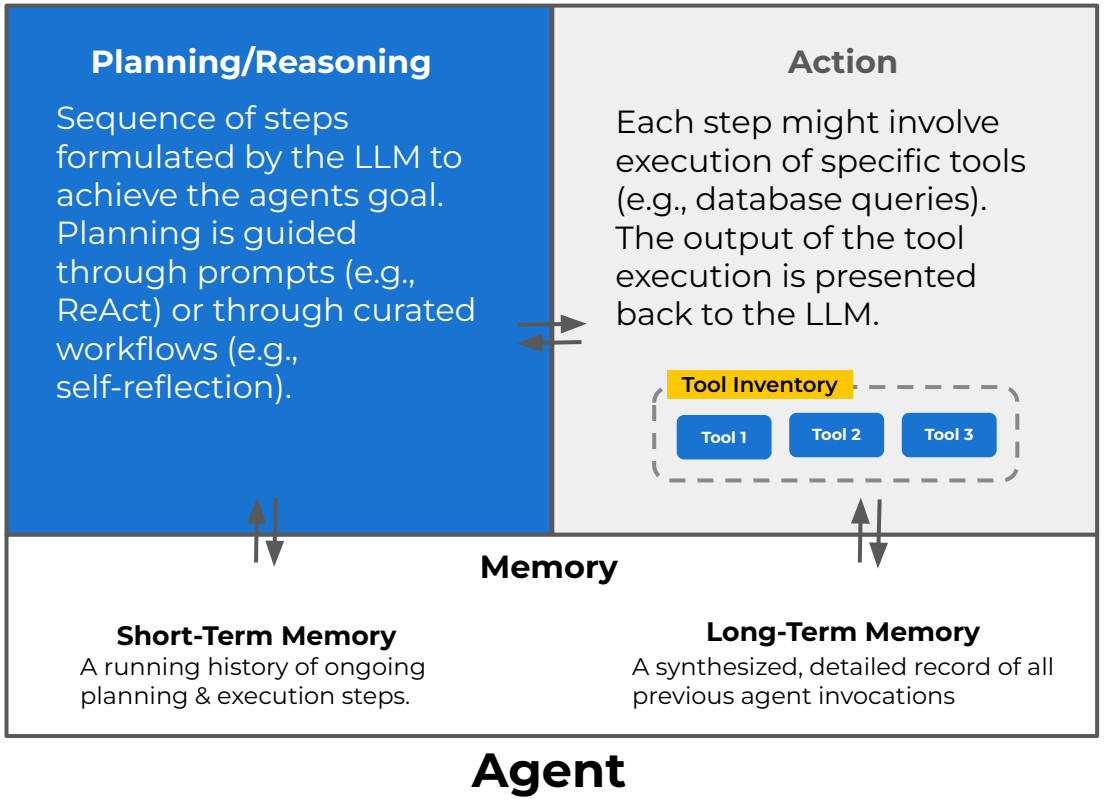

- A single-agent system operates through iterative cycles of planning, action, and memory.
- The LLM formulates a sequence of steps (Planning/Reasoning) to achieve a goal, guided by prompts or curated workflows. Each step may involve the LLM selecting and executing specific tools (Action).
- The results of these tool executions are then fed back into the LLM.
- The agent utilizes both short-term memory to track the current task and long-term memory to synthesize information from previous interactions.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Action in Agents: Tools

Actions can be orchestrated by allowing agents to access tools

Tools are implemented as function calls, with the function's internal logic representing the business task it performs. For example, the Get Order Details tool might be implemented as a get_order_details Python function that queries a database for order status.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

The Plan - Action Interface

The challenge of Inter-operable Agent Tooling

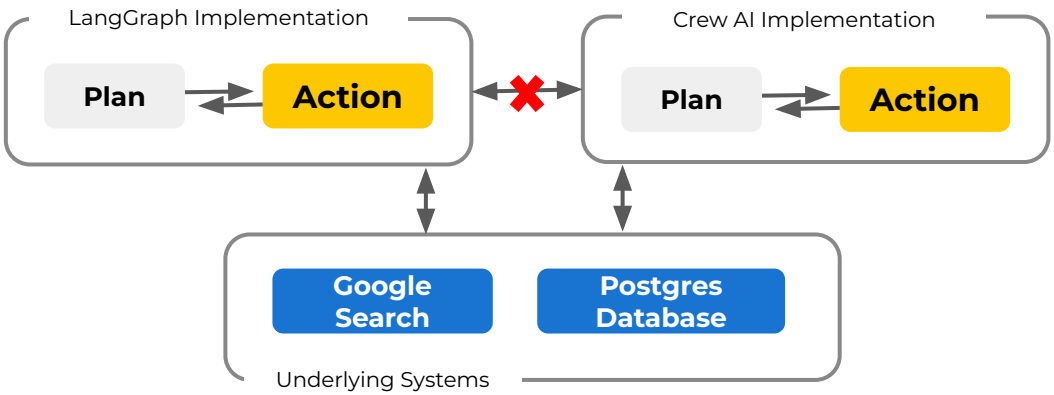

- Planning and Action modules are essential to an agent's functionality, involving frequent back-and-forth interactions.
- However, tool execution in the Action module and its integration with Planning are often implementation-specific.
- Even though the same tools are used by an agent (for e.g., Google Search and Database), implementation specific details don't allow the plan - action interface from one framework to be used with another framework.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

The Plan - Action Interface

Solving Interface Incompatibility with **Model Context Protocol (MCP)**

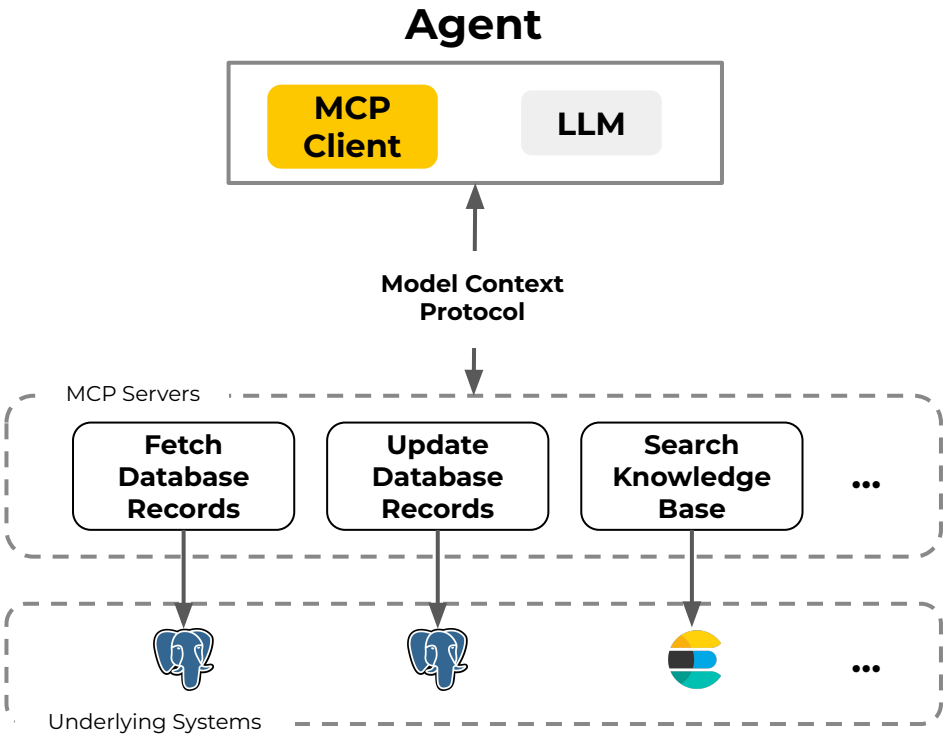

- The MCP framework connects agents to external capabilities using a client-server protocol.
- Through this protocol, MCP servers offer tools (which can modify state or connect to external systems), prompts (reusable templates that guide tool usage), and resources (like data and content).
- On the client side, agents use LLMs to plan which tool to invoke and format the necessary parameters for the request.
- MCP servers securely manage these requests, acting as an interface to underlying systems (such as databases) and replying to the client in a standardized format.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Where Do AI Agents Show Up?

AI agents are already working behind the scenes in everyday tools and services.

Customer Support

Resolve queries, process refunds, and escalate issues — instantly, without human wait times.

Smart Assistants

Set reminders, control devices, manage schedules, and handle daily tasks on command.

Shopping Assistants

Find products, compare prices, and complete purchases tailored to your preferences.

Banking Assistants

Track spending, flag fraud, and automate transfers — your personal financial co-pilot.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Made with GAMMA

Why AI Agents Are a Big Deal

⏱️ Save Time

Automate repetitive work so teams focus on high-value decisions.

⚡ Faster Decisions

Process data and act in seconds, not hours.

📈 Boost Productivity

Handle multiple tasks simultaneously, at scale.

🤝 Always Available

No breaks, no delays — support 24/7, 365 days a year.

The Future of Work

The most powerful organisations won't choose between humans and AI agents — they'll combine both.

Humans bring creativity and judgment. Agents bring speed and scale. Together, they unlock a new level of performance.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Made with GAMMA

Generative AI Agents

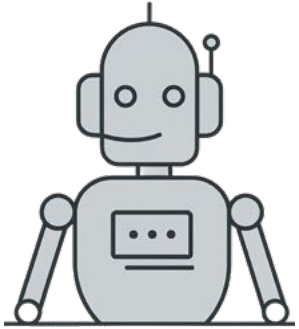

Generative AI agents are an evolving subset of AI Agents that use generative models to perceive their environment, produce novel and relevant outputs, make decisions, and interact with their surroundings to achieve goals with minimal human intervention.

Autonomy – Do not rely on pre-programmed rules or step-by-step instructions and can independently handle complex scenarios relying on reasoning afforded by LLMs.

Perception – Use text-based or multimodal inputs and interactions with other systems or users to perceive their environment.

Action – Determine the best course of action and execute tasks to fulfil their goals. Task execution is achieved by selecting, sequencing and calling a set of tools available to the agent.

Goal-Oriented Behavior – Designed to achieve specific, pre-defined goals by autonomously breaking down complex tasks into smaller steps.

**For the rest of this course, we use the terms agents and Generative AI agents interchangeably*

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Generative AI Agents

Agent = LLM (for planning/reasoning) + Tool Execution (for action)

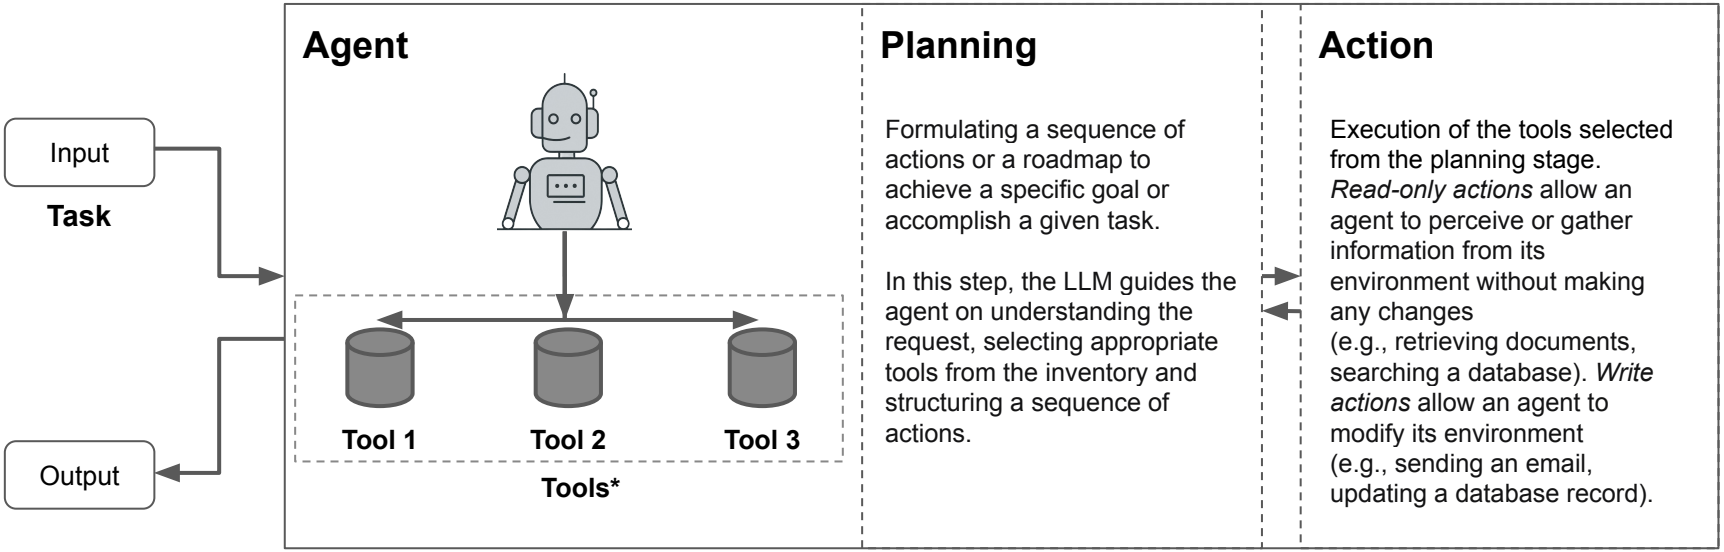

Tools are the specific resources (e.g., APIs, databases) the agent uses to find information or take action beyond its own knowledge. The agent evaluates the output from each tool step, replanning or escalating upon failure. Once successful, it uses the gathered information or completed actions to generate a natural language customer response.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Generative AI Agents: Use Case

Handling E-Commerce Customer Inquiries

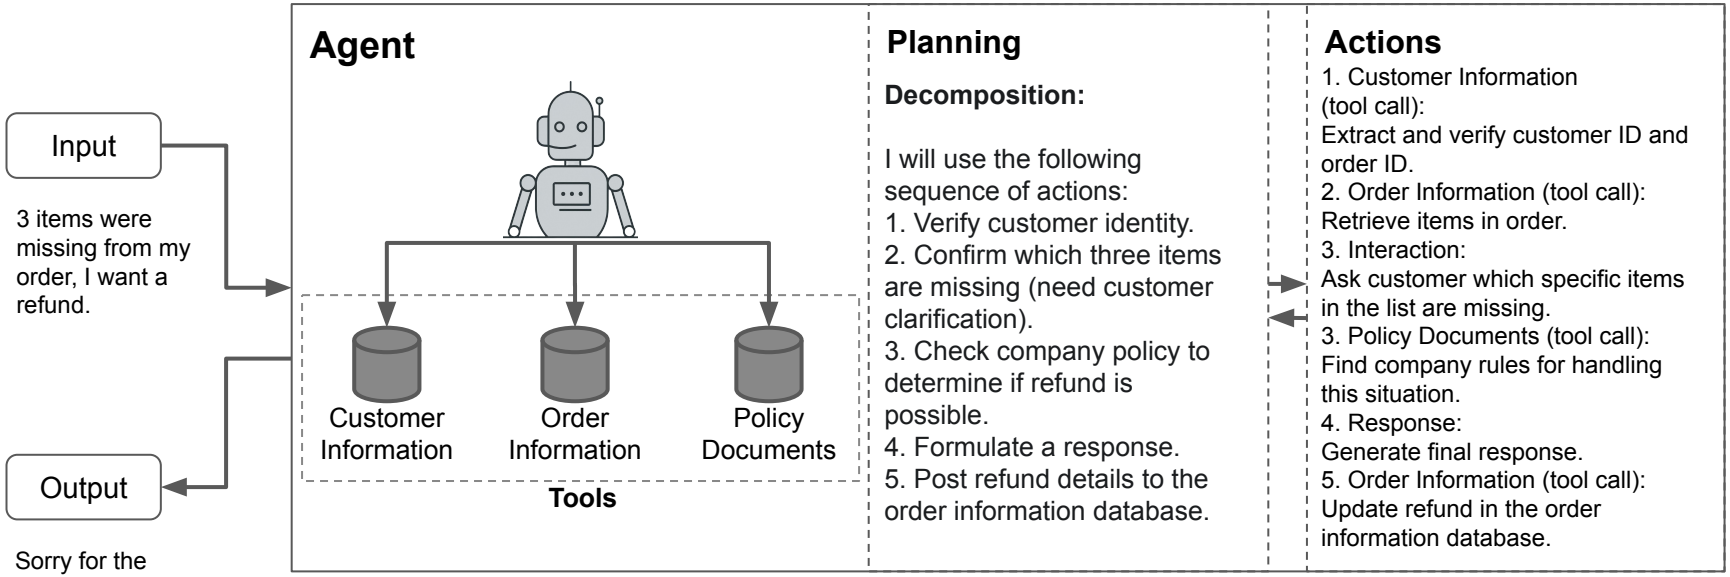

The agent looks at the entire trace of actions and their outputs and composes the final answer. If resolution is not possible within the action sequence, the query is escalated to a human.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Contextualizing GenAI Agents

	LLMs	GenAI Assistants	GenAI Workflows	GenAI Agents
What?	Advanced Machine Learning models that can perform singular natural language tasks.	LLMs that take on specific roles in an interactive manner. Maintain history of previous interactions.	Tasks are broken down into a fixed sequence of steps. Each step is executed by an LLM call (and possibly, tools).	Autonomously break down complex tasks into subtasks and carry them out using the tools at their disposal. Goal-oriented systems.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Contextualizing GenAI Agents

	LLMs	GenAI Assistants	GenAI Workflows	GenAI Agents
When to use?	Excel in tasks like summarization, text generation, translation, and creative writing. To provide coding assistance. Foundational block for applications involving reasoning.	Useful for customer support and service automation. Leverage external data sources for context using RAG. Foundational block for multi-step GenAI workflows.	Ideal for situations where the task can be easily and cleanly decomposed into fixed subtasks (e.g., generating marketing copy). Routing and parallelization can increase effectiveness of independent subtasks (e.g., specialized processes for customer queries).	Ideal for situations where it's difficult or impossible to predict the required number of steps and where a fixed path cannot be hardcoded. Goal-based agents excel in complex decision-making environments, evaluating potential actions, such as in financial trading.

For many organizations, starting with prescriptive GenAI workflows provides a more manageable and often more effective entry point than immediately deploying autonomous agents.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Internal Workings of A Single-agent System

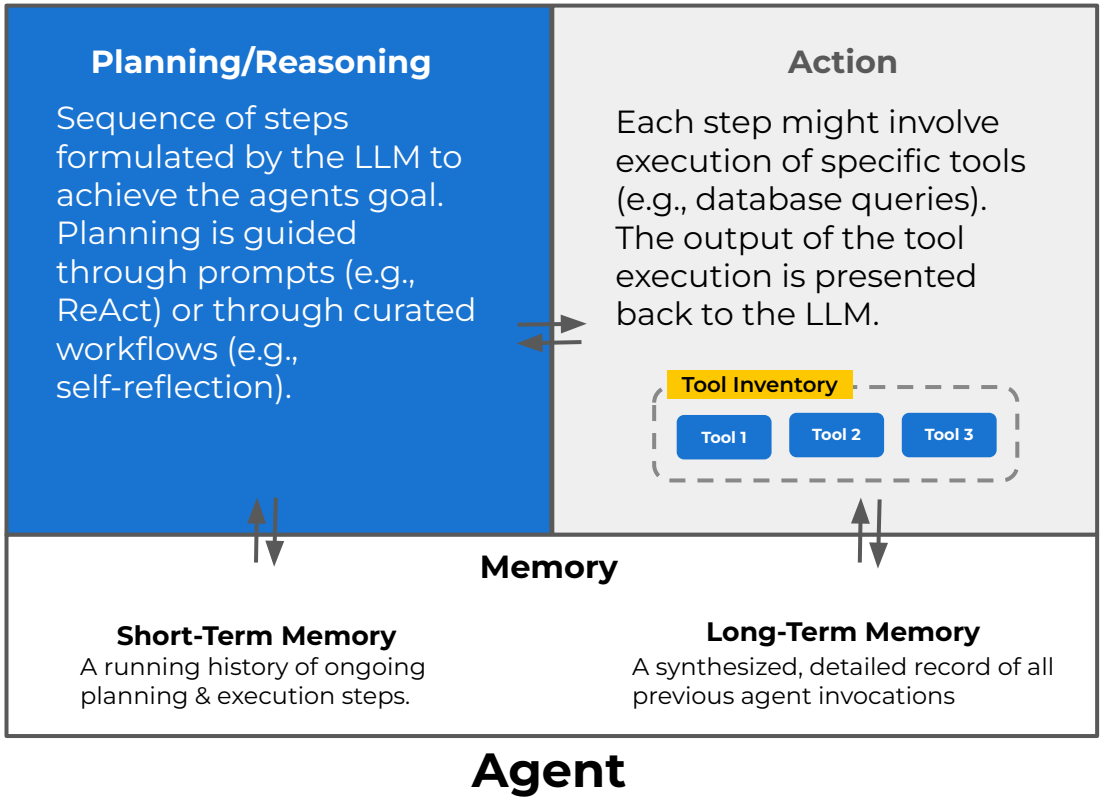

- A single-agent system operates through iterative cycles of planning, action, and memory.
- The LLM formulates a sequence of steps (Planning/Reasoning) to achieve a goal, guided by prompts or curated workflows. Each step may involve the LLM selecting and executing specific tools (Action).
- The results of these tool executions are then fed back into the LLM.
- The agent utilizes both short-term memory to track the current task and long-term memory to synthesize information from previous interactions.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Action in Agents: Tools

Actions can be orchestrated by allowing agents to access tools

Tools are implemented as function calls, with the function's internal logic representing the business task it performs. For example, the Get Order Details tool might be implemented as a get_order_details Python function that queries a database for order status.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

The Plan - Action Interface

The challenge of Inter-operable Agent Tooling

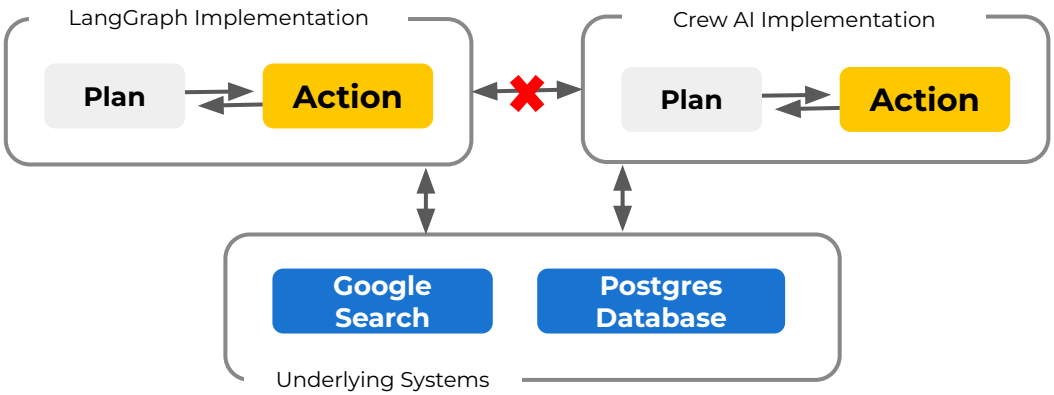

- Planning and Action modules are essential to an agent's functionality, involving frequent back-and-forth interactions.
- However, tool execution in the Action module and its integration with Planning are often implementation-specific.
- Even though the same tools are used by an agent (for e.g., Google Search and Database), implementation specific details don't allow the plan - action interface from one framework to be used with another framework.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

The Plan - Action Interface

Solving Interface Incompatibility with **Model Context Protocol (MCP)**

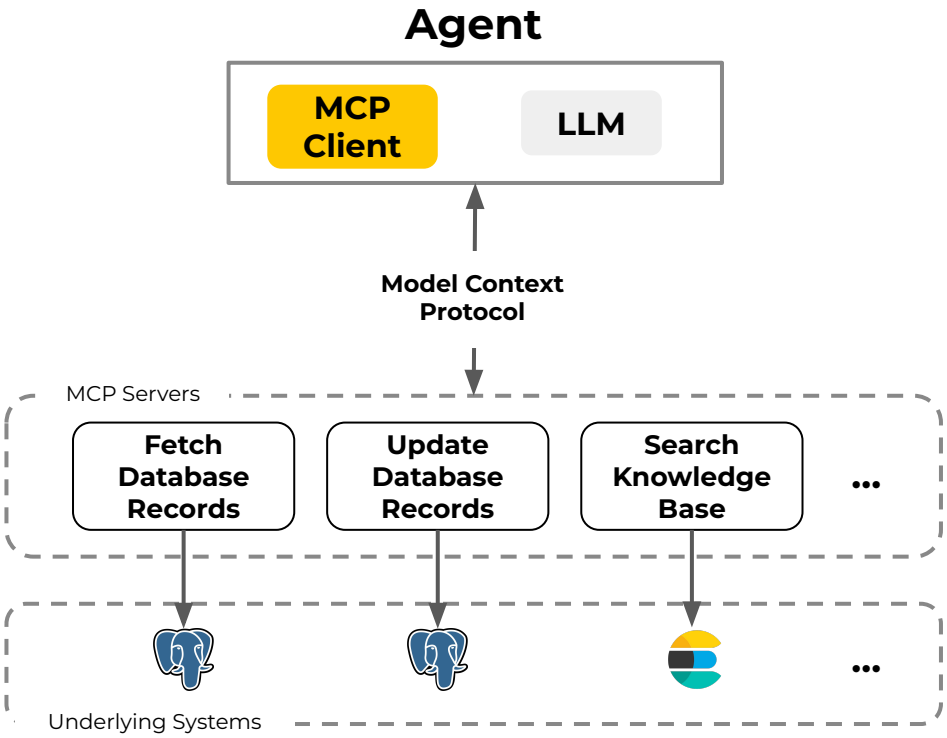

- The MCP framework connects agents to external capabilities using a client-server protocol.
- Through this protocol, MCP servers offer tools (which can modify state or connect to external systems), prompts (reusable templates that guide tool usage), and resources (like data and content).
- On the client side, agents use LLMs to plan which tool to invoke and format the necessary parameters for the request.
- MCP servers securely manage these requests, acting as an interface to underlying systems (such as databases) and replying to the client in a standardized format.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Agent Planning: ReAct = Reason + Act

Internal reasoning can be guided by prompts

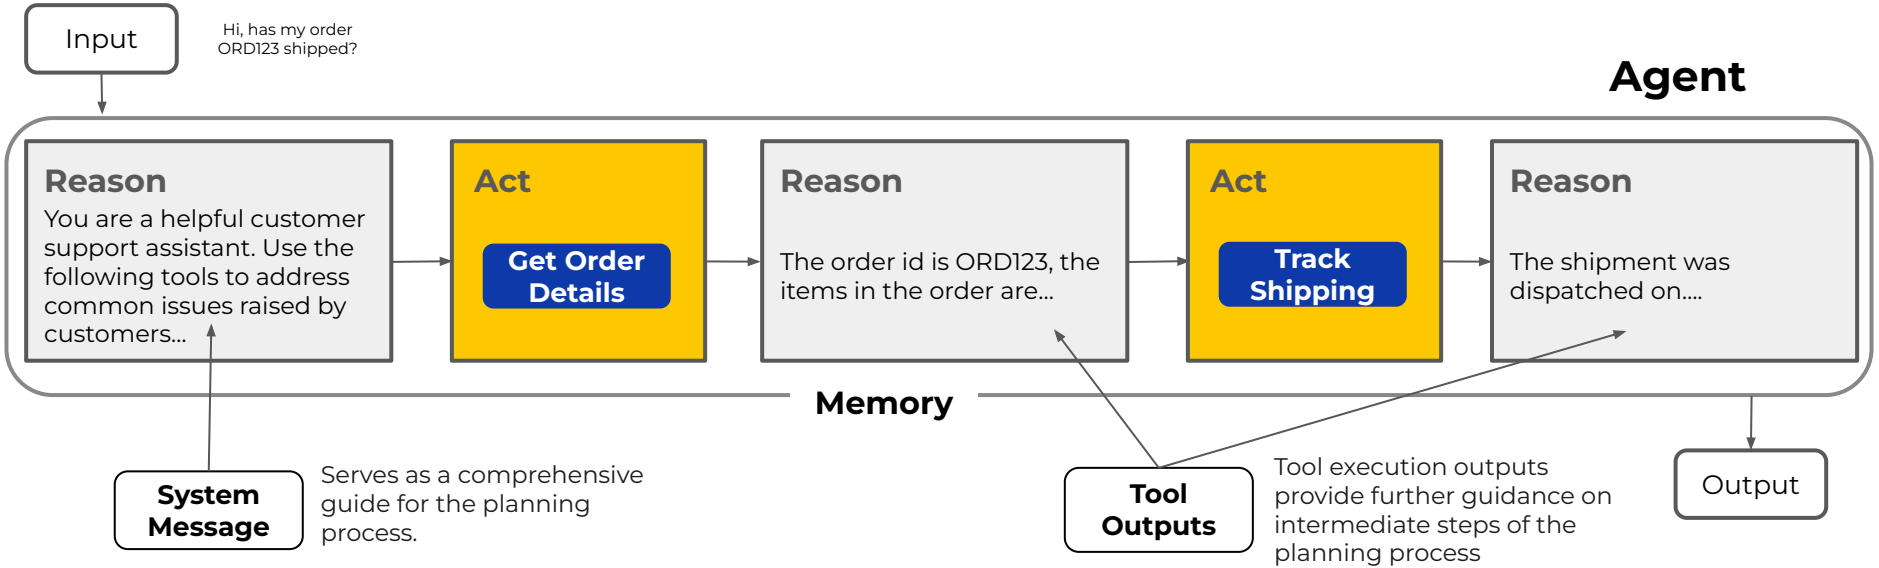

In ReAct, the agent alternates between Reasoning (planning which tool to use, guided by system messages and previous steps) and Acting (executing that tool). Tool results feed back into the agent's memory, informing subsequent reasoning and actions until the final objective is reached and an output is generated.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Agent Planning: Self-Reflection

Fine-grained internal reasoning can be achieved using workflows

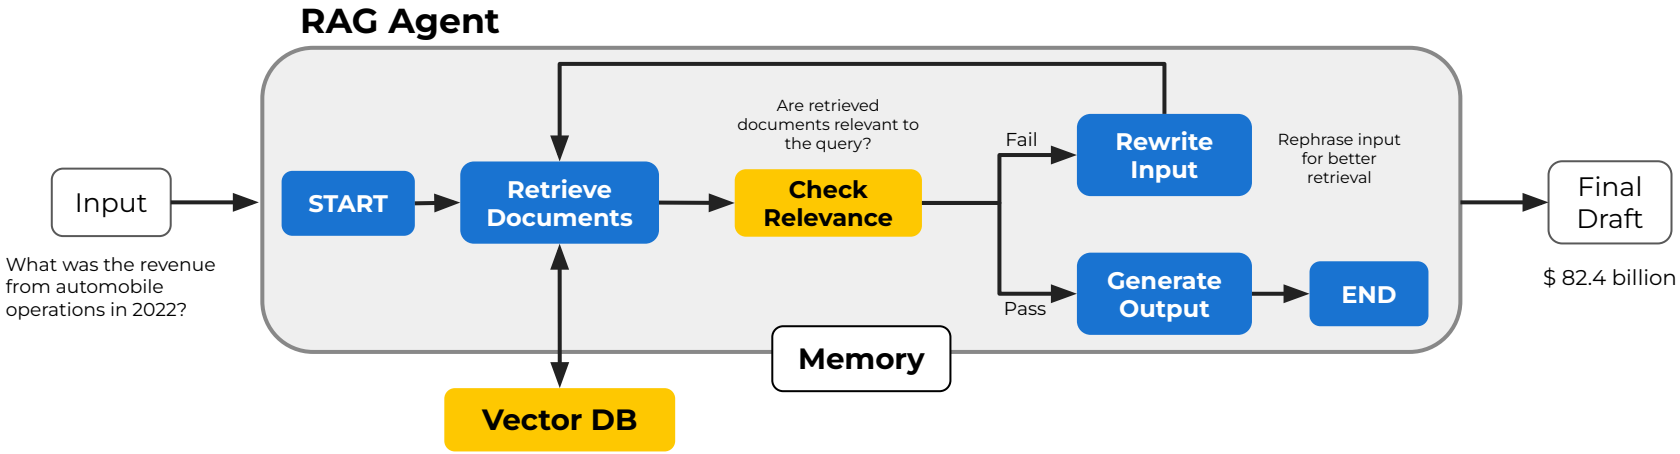

- **Structured Reasoning:** Agent operations are defined as an explicit workflow, breaking down tasks into sequential steps like retrieving documents, checking relevance, and generating output.
- **Internal Checkpoints & Correction:** The workflow includes decision points (e.g., "Check Relevance") that assess intermediate results; failure triggers a corrective sub-process (like "Rewrite Input") before retrying.
- **Resource Integration:** External resources like Vector DBs and internal memory are integrated at specific stages within the workflow to support retrieval and maintain context during the process.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Agent Planning: Reflection

Agents can be embedded into workflows for external validation

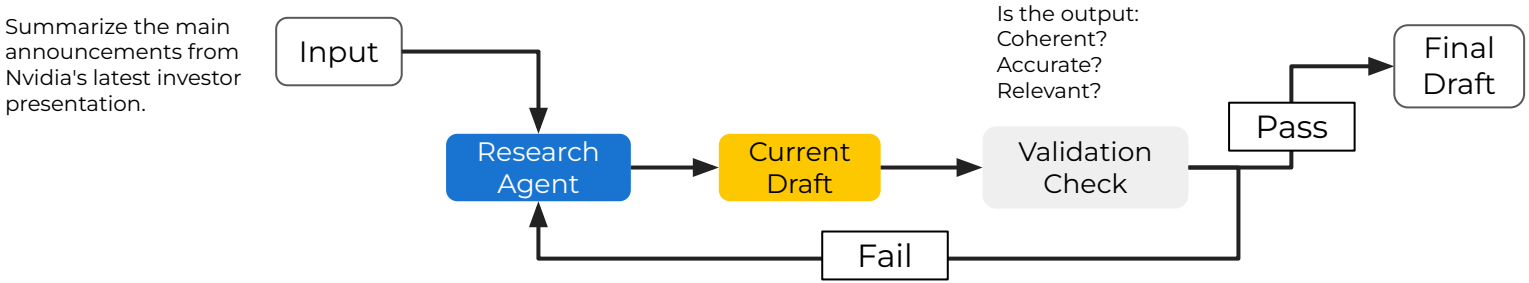

- **Generate-Evaluate-Refine Loop:** In this version, the agent is embedded in a cycle where an agent produces an output, and this output is evaluated against criteria, and the agent revises its work if deficiencies are found.
- **Feedback Guides Revision:** When validation fails, the generated critique (feedback) isn't just noted; it directly informs the agent's next attempt, prompting specific improvements.
- **Triggers Deeper Re-Processing:** Receiving critical feedback often forces the agent to revisit its reasoning, planning, or even re-engage tools to gather/verify information for a better revision.
- **Progressive Output Enhancement:** This iterative process aims to systematically enhance the quality and reliability of the agent's final output before it's considered complete.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Managing Memory

Using Layered Memory Improves Agent Performance

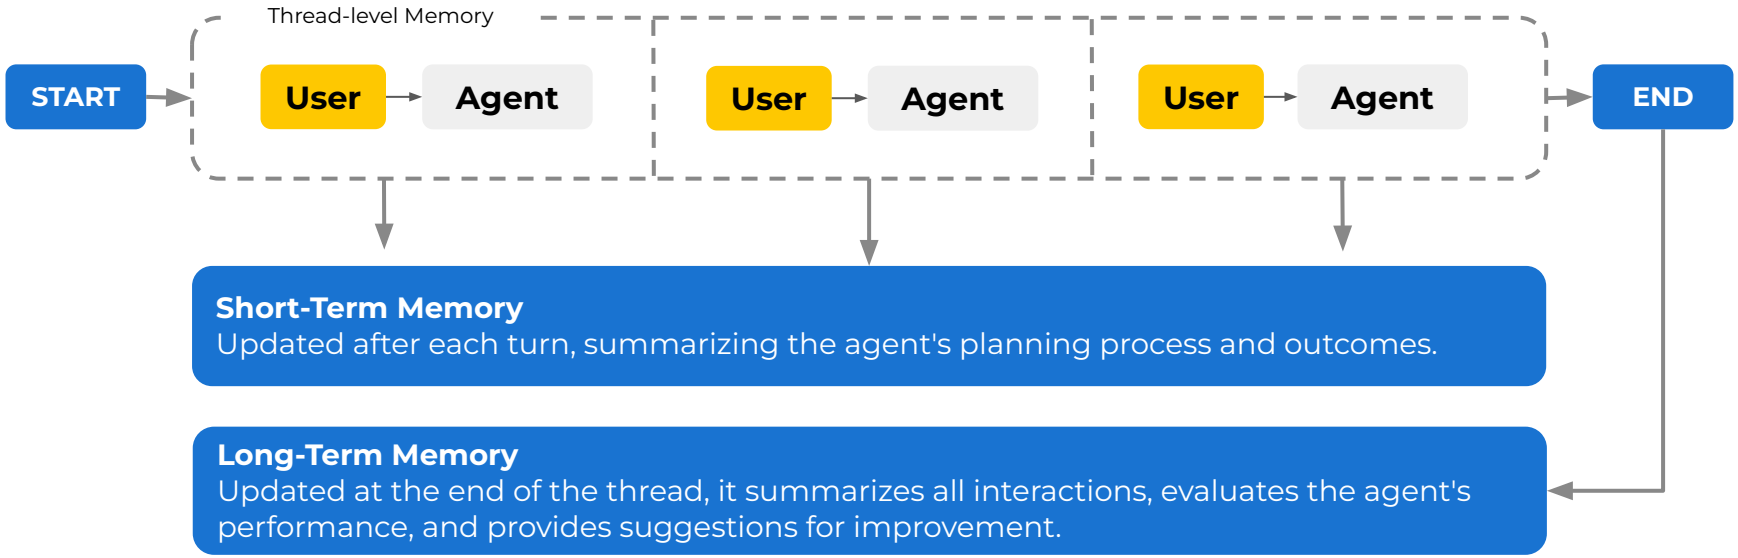

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Self-Reflection

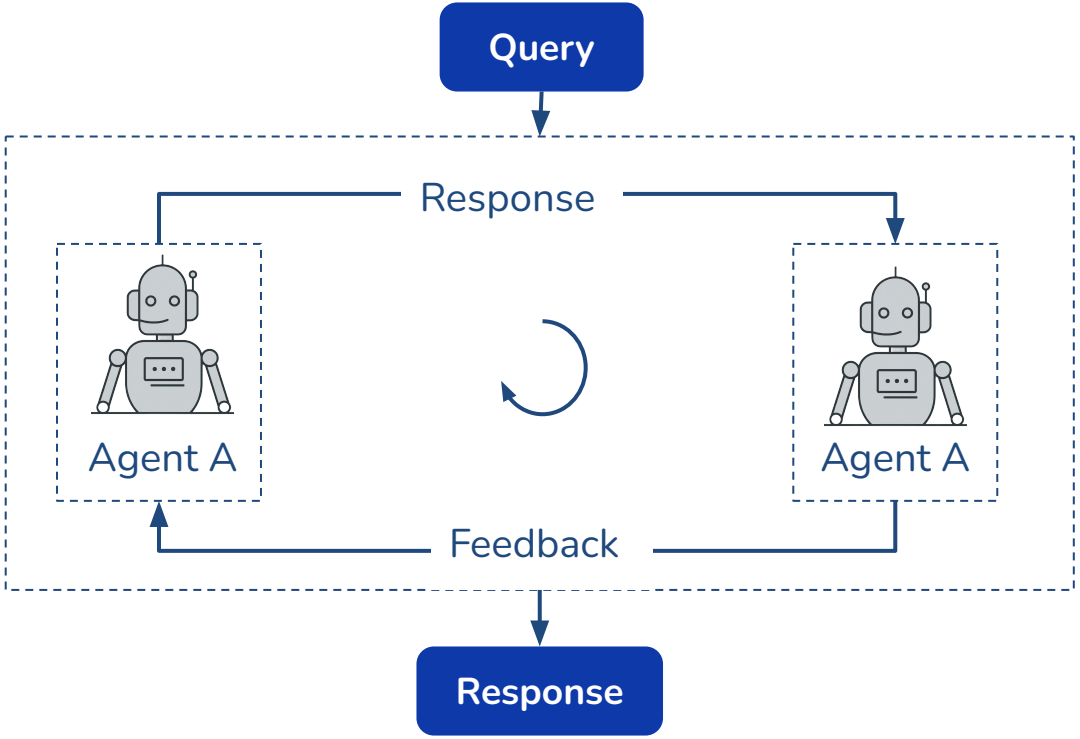

Self-reflective agents iteratively refine their own answer based on an assessment of: adherence to instructions and quality of response.

Execution concludes either when a fixed number of iterations are done or when no further improvement is possible.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Demonstration

Hands-on: Demonstration using n8n

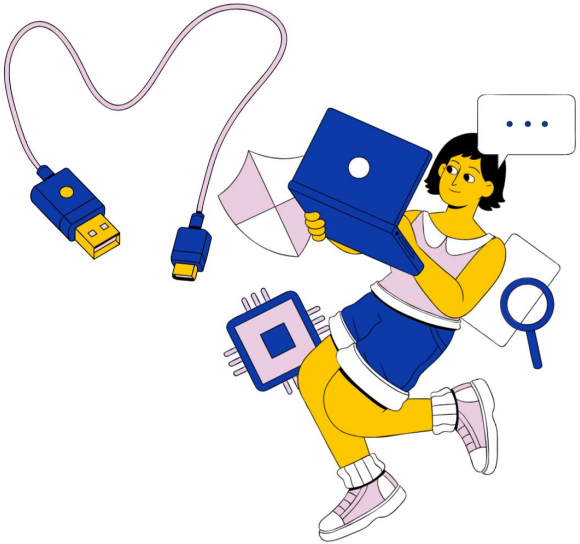

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Why Humans Remain Crucial

Why Humans Remain Crucial

The AI Revolution

AI agents process vast datasets, identify complex patterns, and automate intricate tasks at speeds impossible for humans. This unlocks unprecedented efficiency and innovation across every sector.

The Bottleneck

Fully autonomous AI can generate errors, produce "hallucinations," and lack nuanced judgment—especially in high-stakes decisions. This creates trust, reputational, and ROI risks.

The Solution: HITL

Strategic integration of human oversight into AI workflows combines AI's speed with human judgment for reliable, trustworthy, and scalable systems. It's about augmenting—not replacing—expertise.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Questions?

Happy Learning!